

Введение в нейронные сети. Урок 2. Keras



GeekBrains

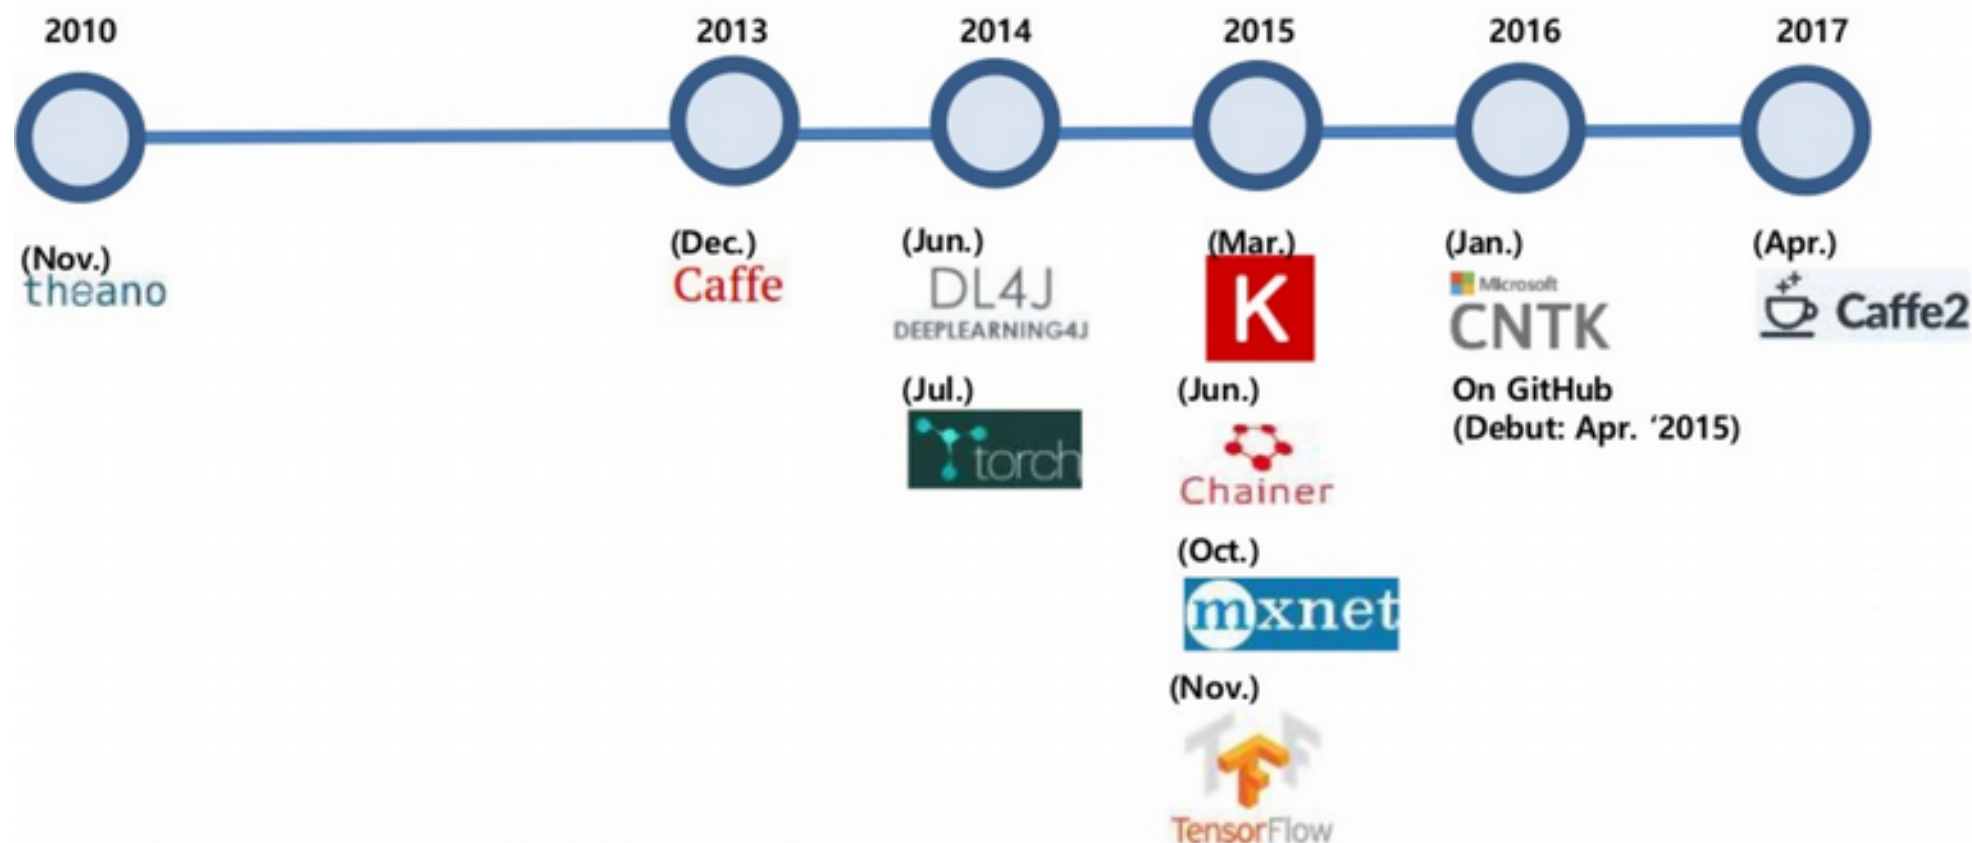
План вебинара



1. Инструменты для создания нейронных сетей.
2. Общие сведения о Keras
3. Синтаксис Keras
4. Практика



Инструменты для создания нейр. сетей



Общие сведения о Keras



ОСНОВЫ СИНТАКСИСА



GeekBrains

Deep Learning with Keras :: CHEAT SHEET



Intro

Keras is a high-level neural networks API developed with a focus on enabling fast experimentation. It supports multiple backends, including TensorFlow, CNTK and Theano.

TensorFlow is a lower level mathematical library for building deep neural network architectures. The Keras R package makes it easy to use Keras and TensorFlow in R.



<https://keras.rstudio.com>

<https://www.manning.com/books/deep-learning-with-r>

INSTALLATION

The Keras R package uses the Python keras library. You can install all the prerequisites directly from R. https://keras.rstudio.com/reference/install_keras.html

```
library(keras)
install_keras()
```

See `keras_install` for GPU instructions

This installs the required libraries in an Anaconda environment or virtual environment 'r-tensorflow'.

Working with keras models

DEFINE A MODEL

`keras_model()` Keras Model

`keras_model_sequential()` Keras Model composed of a linear stack of layers

`multi_gpu_model()` Replicates a model on different GPUs

COMPILE A MODEL

`compile(object, optimizer, loss, metrics = NULL)` Configure a Keras model for training

FIT A MODEL

`fit(object, x = NULL, y = NULL, batch_size = NULL, epochs = 10, verbose = 1, callbacks = NULL, ...)` Train a Keras model for a fixed number of epochs (iterations)

`fit_generator()` Fits the model on data yielded batch-by-batch by a generator

`train_on_batch()` `test_on_batch()` Single gradient update or model evaluation over one batch of samples

EVALUATE A MODEL

`evaluate(object, x = NULL, y = NULL, batch_size = NULL)` Evaluate a Keras model

`evaluate_generator()` Evaluates the model on a data generator

PREDICT

`predict()` Generate predictions from a Keras model

`predict_proba()` and `predict_classes()` Generates probability or class probability predictions for the input samples

`predict_on_batch()` Returns predictions for a single batch of samples

`predict_generator()` Generates predictions for the input samples from a data generator

OTHER MODEL OPERATIONS

`summary()` Print a summary of a Keras model

`export_savedmodel()` Export a saved model

`get_layer()` Retrieves a layer based on either its name (unique) or index

`pop_layer()` Remove the last layer in a model

`save_model_hdf5()` `load_model_hdf5()` Save/Load models using HDF5 files

`serialize_model()` `unserialize_model()` Serialize a model to an R object

`clone_model()` Clone a model instance

`freeze_weights()` `unfreeze_weights()` Freeze and unfreeze weights

CORE LAYERS

`layer_input()` Input layer

`layer_dense()` Add a densely-connected NN layer to an output

`layer_activation()` Apply an activation function to an output

`layer_dropout()` Applies Dropout to the input

`layer_reshape()` Reshapes an output to a certain shape

`layer_permute()` Permute the dimensions of an input according to a given pattern

`layer_repeat_vector()` Repeats the input n times

`layer_lambda(object, f)` Wraps arbitrary expression as a layer

`layer_activity_regularization()` Layer that applies an update to the cost function based input activity

`layer_masking()` Masks a sequence by using a mask value to skip timesteps

`layer_flatten()` Flattens an input

TRAINING AN IMAGE RECOGNIZER ON MNIST DATA

```
# input layer: use MNIST images
mnist <- dataset_mnist()
x_train <- mnist$train$x; y_train <- mnist$train$y
x_test <- mnist$test$x; y_test <- mnist$test$y
```

```
# reshape and rescale
x_train <- array_reshape(x_train, c(nrow(x_train), 784))
x_test <- array_reshape(x_test, c(nrow(x_test), 784))
x_train <- x_train / 255; x_test <- x_test / 255
```

```
y_train <- to_categorical(y_train, 10)
y_test <- to_categorical(y_test, 10)
```

```
# defining the model and layers
model <- keras_model_sequential()
model %>%
  layer_dense(units = 256, activation = 'relu',
    input_shape = c(784)) %>%
  layer_dropout(rate = 0.4) %>%
  layer_dense(units = 128, activation = 'relu') %>%
  layer_dense(units = 10, activation = 'softmax')
```

```
# compile (define loss and optimizer)
model %>% compile(
  loss = 'categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy'))
```

```
# train (fit)
model %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 128,
  validation_split = 0.2)
model %>% evaluate(x_test, y_test)
model %>% predict_classes(x_test)
```



RStudio® is a trademark of RStudio, Inc. • CC BY SA RStudio • <https://www.rstudio.com> • 844-448-1212 • info@rstudio.com • Learn more at keras.rstudio.com • Keras 2.1.2 • Updated: 2017-12

Структура Keras

Models

- Sequential
- Model API


Layers

- сверточные
- рекуррентные
- полносвязные
- служебные

Preprocessing

- utils
- обработка изображений
- обработка текстов

Модель



Models.Model

`keras.Model()`

`inputs`
`outputs`
`name`

`Model.summary()`

Имя

Слой 1 $f(s)$

ней ней нейрон

Слой 2


Слой 3

Models.Sequential

`keras.Sequential()`

`layers,`
`Name`

`модель.add(<Слой>)`

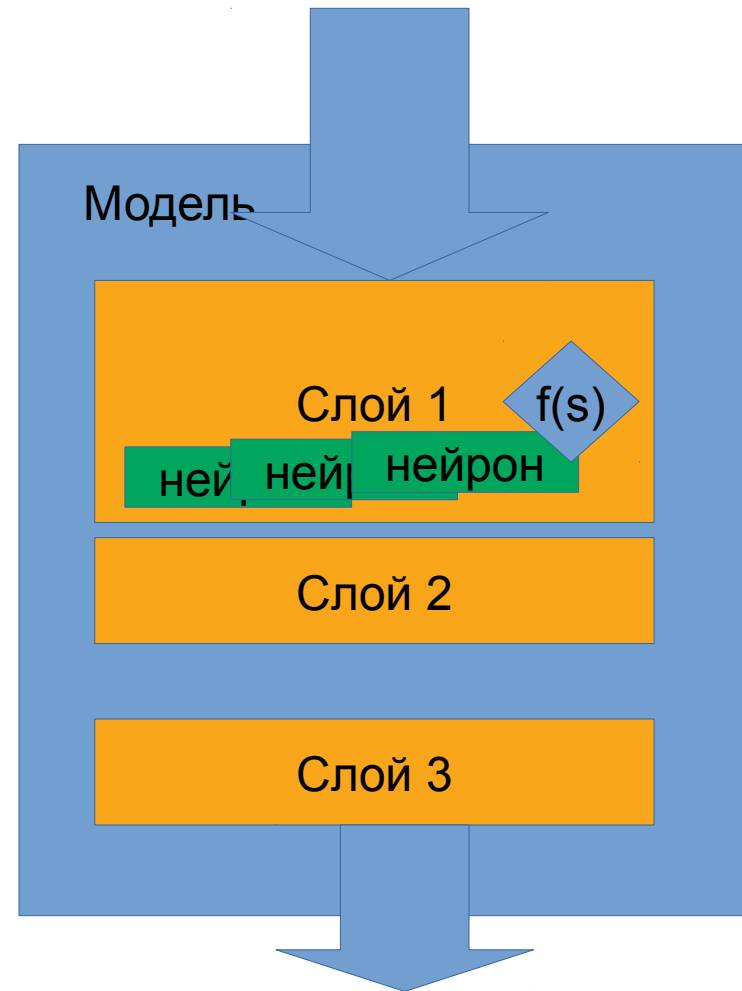


Models Model training API

```
Model.compile(  
    optimizer="rmsprop",  
    loss=None,  
    metrics=None)
```

```
Model.fit( x=None, y=None,  
    batch_size=None,  
    epochs=1,  
    validation_split=0.0)
```

```
Model.predict( x)
```



Вопросы



GeekBrains



Практическое задание



1. Попробуйте обучить нейронную сеть на Keras на Fashion-MNIST датасете. Опишите в комментарии к уроку - какой результат вы добились от нейросети? Что помогло вам улучшить ее точность?
- *2. Поработайте с документацией Keras. Найдите полезные команды не разобранные на уроке.